

**Duke University**

Department of Mechanical Engineering & Materials Science (MEMS)

# Final Project

RoboBarista

*An Autonomous Robotic Barista System*

*Submitted by*

## Team 7

Alexey Khotimsky (Ph.D. in MEMS)

Mohammad Afrazi (Ph.D. in MEMS)

Celia Pan (M.S. in MEMS)

Yongjae Lee (M.S. in MEMS)

Zhe Feng (M.S. in MEMS)

**Course:** ECE 383 / ME 555 – Introduction to Robotics

**Instructor:** Dr. Siobhan Rigby Oca

December 08, 2025

# Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Documentation of Robot Actions</b>	<b>3</b>
2.1	System Architecture	3
2.2	Visual Perception and Gesture Recognition	3
2.2.1	Model Selection and Dataset	3
2.2.2	Training Configuration	4
2.2.3	Inference and ROI Optimization	4
2.3	Granular Description of Efforts	4
2.4	Robot Execution Image	5
2.5	Code Explanation	5
2.5.1	YOLO Detection Node (Perception Layer)	5
2.5.2	GestureControl (Supervisor Node)	6
2.5.3	Motion Bots (Execution Layer)	6
2.5.4	Digital Scale Node	6
2.6	End-Effector Cup Design for Stable Grasping	7
<b>3</b>	<b>Human Integration</b>	<b>7</b>
<b>4</b>	<b>Challenges</b>	<b>8</b>
<b>5</b>	<b>Future Work</b>	<b>8</b>
<b>6</b>	<b>Takeaways</b>	<b>8</b>
<b>A</b>	<b>Code Appendix</b>	<b>11</b>
A.1	Gesture Detection Node	11
A.2	Gesture Control (Supervisor Node)	16
A.3	Simple Pour Node (Static)	19
A.4	Barista Pour Node (Parametric)	24
A.5	Arduino Code for Scale	28
A.6	YOLO Dataset Configuration (data.yaml)	30
A.7	Computer Vision: Training Script	30

## List of Figures

1	RoboBarista at Duke! (The image created by ChatGPT and edited by the team.)	2
2	System-level view of RoboBarista. Left: System diagram showing data flow: Camera → YOLO → Gesture Control → MoveIt2. Right: runtime screenshot showing ROI hand detection, RViz visualization, and ROS 2 console logs that confirm the gesture-triggered pouring sequence.	3
3	RoboBarista hardware setup. Left: the Kinova Gen3 Lite performing the parametric circular pouring task. Right: the custom digital scale with a load cell and Arduino-based electronics used for closed-loop gravimetric pouring.	5
4	Custom cup designs and prototypes. Top: CAD models of the smooth, ribbed, and diamond patterned cups. Bottom: 3D printed prototypes used to evaluate grasp stability. The diamond patterned cup provided the most robust grasp and was selected for all final trials.	7
5	Flowchart of the Human-Robot Interaction Logic.	9

## 1 Introduction

The deployment of service robots in unstructured environments requires capabilities beyond simple pick-and-place operations. Specifically, tasks involving the manipulation of fluids or granular materials demand precise, dynamic trajectory control to ensure spill-free transport and even distribution. This project explores the development of an autonomous “Robotic Barista” system designed to bridge the gap between static automation and dynamic human service tasks.

Our system utilizes a **Kinova Gen3 Lite** manipulator driven by a **ROS 2** (Robot Operating System) architecture. The primary objective was to implement a closed-loop control system that allows a human operator to command specific pouring profiles using natural hand gestures. By integrating computer vision with real-time motion planning, we demonstrate a robust framework for Human-Robot Collaboration (HRC).

The core technical contribution of this work is the development of a **parametric trajectory generator** for the pouring motion. Rather than relying on static waypoints, the system calculates a continuous circular path—mimicking the technique of a professional barista—to ensure uniform extraction and distribution of the coffee grounds. This report details the kinematic derivation, software architecture, and experimental validation of the system.

In our physical experiments, we use uncooked rice as a granular stand-in for coffee to simplify cleanup and enable precise gravimetric measurements on the scale.

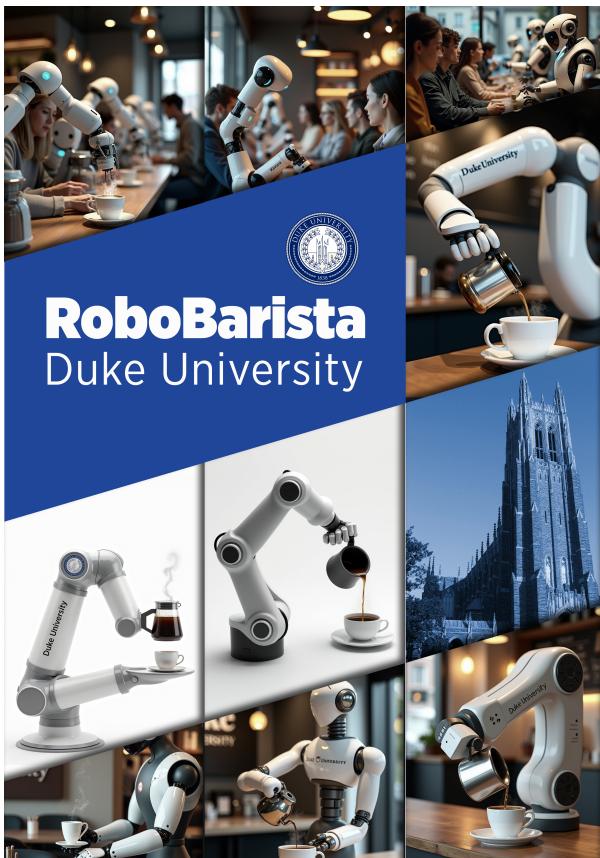


Figure 1: RoboBarista at Duke! (The image created by ChatGPT and edited by the team.)

## 2 Documentation of Robot Actions

### 2.1 System Architecture

The system integration is illustrated in Figure 2. The communication pipeline flows from the vision sensors to the central supervisor node, which dispatches motion tasks to the manipulator.

In the perception layer, a camera stream is processed by a YOLOv8-based gesture recognizer that publishes discrete commands on the `/gesture` topic. The `GestureControl` node subscribes to this topic and selects either the *Simple Gravimetric Pour* or the *Barista Style Circular Pour* behavior node. Both pouring nodes use MoveIt2 to plan and execute Cartesian motions for the Kinova Gen3 Lite arm.

For the gravimetric pour, we additionally integrate a digital scale composed of a load cell, an amplifier board, and an Arduino microcontroller. The Arduino streams the measured weight over a serial connection to a ROS 2 node, which publishes the real-time mass on a `/weight` topic. The SimplePourBot monitors this topic during the slow tilting motion and stops pouring as soon as the measured mass crosses a 100 g threshold, then returns the cup to an upright pose. This closes the loop around the poured mass, while the barista mode remains open loop and focuses on generating a smooth circular trajectory without using the weight feedback.

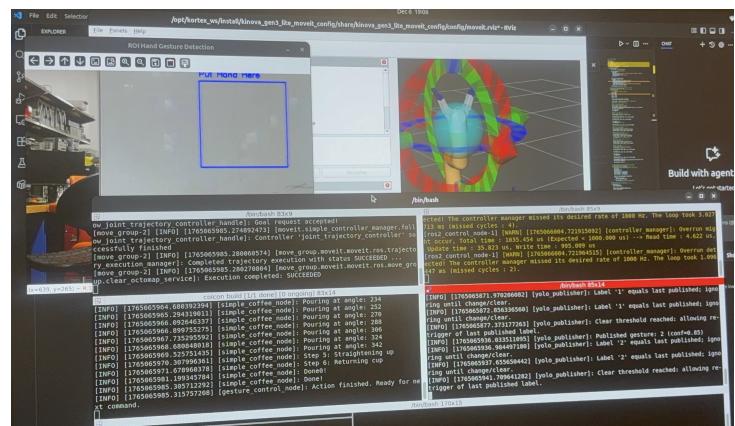


Figure 2: System-level view of RoboBarista.

Left: System diagram showing data flow: Camera → YOLO → Gesture Control → MoveIt2.  
Right: runtime screenshot showing ROI hand detection, RViz visualization, and ROS 2 console logs that confirm the gesture-triggered pouring sequence.

### 2.2 Visual Perception and Gesture Recognition

#### 2.2.1 Model Selection and Dataset

For the gesture recognition task, we selected the YOLOv8n architecture. This model was chosen for its lightweight footprint and high inference speed, which is essential for real-time robotic interaction without significant latency.

We utilized an open-source Hand Gesture Recognition Dataset sourced from Open Source Toolkit<sup>1</sup>, which contains labeled images for numerical gestures ranging from 0 to 9. The model was trained on the full dataset to ensure robust feature extraction and generalization capability across all hand signs. For this specific application, we implemented a filtering logic in the ROS node, where the system exclusively monitors for classes “1” and “2”, mapping them to the robot’s “simple Pour (Simple Gravimetric Pour)” and “Barista Pour (Barista Style Circular Pour)” modes respectively.

<sup>1</sup>Dataset available at: <https://gitcode.com/open-source-toolkit/2c695>

### 2.2.2 Training Configuration

The model was trained using the Ultralytics framework via a transfer learning approach. We initialized the model with pretrained COCO weights (`yolov8n.pt`) to leverage learned feature extractors, significantly reducing the training time required for our custom dataset.

The training parameters were configured as follows:

- Epochs: 30
- Batch Size: 16
- Image Size:  $640 \times 640$  pixels
- Optimizer: SGD (Stochastic Gradient Descent)
- Hardware: Training was performed on a CPU-based environment

### 2.2.3 Inference and ROI Optimization

A key challenge in computer vision for open environments is background noise (e.g., detecting a face or a hand in the background as a command). To mitigate this, we implemented a Region of Interest (ROI) logic in the inference script.

Instead of processing the entire camera frame, the system defines a specific active zone (200x200 pixels) located in the top-right quadrant of the video feed. The workflow is as follows:

- Crop: The script extracts only the ROI area from the main frame.
- Inference: The YOLOv8 model runs detection only on this cropped section.
- Mapping: If a gesture is detected with a confidence score above 0.5, the local coordinates are mapped back to the global frame for visualization.

This approach effectively acts as a safety filter, ensuring the robot only responds when the user deliberately places their hand in the designated "command zone."

## 2.3 Granular Description of Efforts

The robot's operation is divided into four distinct phases:

1. **Initialization Phase:** The robot moves to a safe "Home" configuration, opens the gripper, and waits in an idle state with the cup placed on the table.
2. **Perception & Command Phase:** The vision system monitors the region of interest and publishes a gesture label on the `/gesture` topic. Upon receiving a valid command, the `GestureControl` node selects the corresponding pouring behavior, and the arm computes inverse kinematics to move to a pre-grasp pose and securely grasp the cup.
3. **Dispensing Phase:** After grasping the cup, the system executes one of two pouring behaviors:
  - *Simple Gravimetric Pour (Command "1")*: The arm moves the cup above the bowl and calls a slow-tilt routine that interpolates between the upright and pouring orientations over 20 small quaternion steps. During this motion, the system reads the weight from the load cell and Arduino-based scale in real time. When the measured mass crosses a 100 g threshold, the node stops the tilt, returns the cup to an upright pose, and prepares to place it back on the table.

- *Barista Style Circular Pour (Command “2”)*: The arm tilts the cup once to a fixed angle and then follows a parametric circular trajectory of radius  $r = 3\text{ cm}$  above the bowl. This mode ignores the scale feedback and focuses on producing a smooth, visually pleasing circular pour pattern similar to a manual barista pour.
- 4. Return Phase:** The robot brings the cup back to the table, opens the gripper to release it, and then returns to the “Home” pose, resetting the system for the next gesture command.

## 2.4 Robot Execution Image

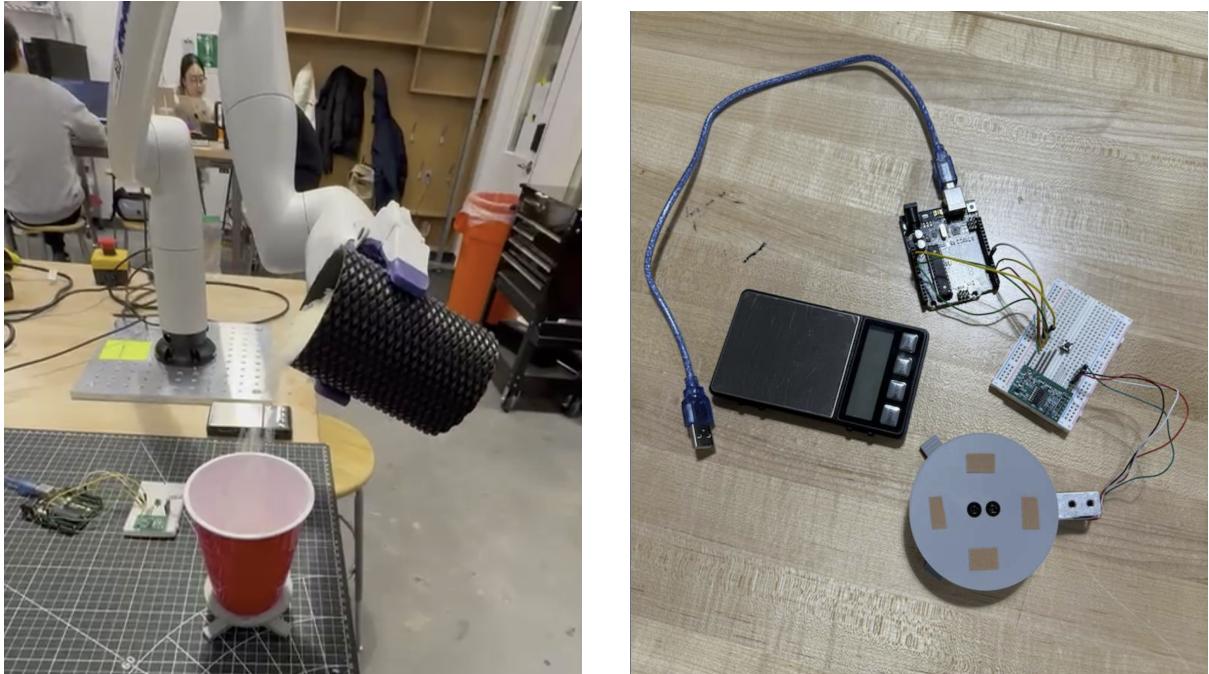


Figure 3: RoboBarista hardware setup.

Left: the Kinova Gen3 Lite performing the parametric circular pouring task.

Right: the custom digital scale with a load cell and Arduino-based electronics used for closed-loop gravimetric pouring.

## 2.5 Code Explanation

The software is architected using a modular **Supervisor–Worker** pattern implemented through three ROS 2 components responsible for perception, supervision, and execution. The primary control signal flows through the `/gesture` topic, while an additional `/weight` topic connects the digital scale node to the gravimetric pouring behavior. We detail each component below.

### 2.5.1 YOLO Detection Node (Perception Layer)

The YOLO gesture recognizer is implemented as an independent ROS 2 publisher node (`detect_pub.py`). Rather than invoking robot behaviors directly, it performs real-time inference on the camera feed and publishes a filtered gesture label (e.g., “pour”, “barista”) on the `/gesture` topic using `std_msgs/String`.

To ensure that only intentional and stable human gestures trigger robot motion, the node implements several lightweight filtering mechanisms:

- **Multi-frame debouncing**: A gesture is accepted only after appearing consistently for several frames.

- **Gesture-change rule:** The node stores the last published label and only republishes when a new, different gesture is detected. This prevents repeated activations while the user holds the same gesture steady.
- **Clear-frame reset:** After a sequence of “no-detection” frames, the previous gesture is cleared, allowing the same gesture to be triggered again later.

These mechanisms operate above the ROS 2 publish–subscribe layer and ensure that only stable, meaningful gesture commands propagate to the supervisory control layer.

### 2.5.2 GestureControl (Supervisor Node)

This is the central decision-making node. It subscribes to the `/gesture` topic and acts as a dispatcher. Upon receiving a command string, it selects the appropriate behavior:

- **Input “1” / “pour”:** Triggers the `SimplePourBot` thread for a standard static pour.
- **Input “2” / “barista”:** Triggers the `BaristaPourBot` thread for the advanced circular pouring task.

It utilizes a `MultiThreadedExecutor` to spawn these actions in background threads, ensuring that the vision callbacks are never blocked by robot motion. A busy-lock mechanism ensures that only one motion routine executes at a time, even though all components reside in the same ROS package.

### 2.5.3 Motion Bots (Execution Layer)

Two motion scripts implement the robot’s physical behaviors:

- `coffee_pourer.py`: Implements the closed-loop gravimetric pour. The arm moves the cup above the bowl, then calls a slow tilt routine that interpolates between the upright and pouring orientations over 20 small quaternion steps while monitoring the `/weight` topic.
- `coffee_pourer_v2.py`: Implements the open loop barista-style circular pour, where the cup is tilted once to a fixed angle and then follows a circular trajectory.

The circular pour routine computes a sequence of 20 waypoints according to:

$$x_i = x_c + r \cos(\theta_i), \quad y_i = y_c + r \sin(\theta_i) \quad (1)$$

From the supervisor node’s perspective, both behaviors are encapsulated as worker classes, and the node simply instantiates the appropriate bot class and executes its `run_my_code()` method when triggered.

Overall, this architecture cleanly separates perception, supervisory control, and execution responsibilities. The YOLO node provides stable gesture intent, the GestureControl node handles behavior selection and concurrency, and the motion bots perform the corresponding robotic actions through MoveIt 2 and the gripper interface.

### 2.5.4 Digital Scale Node

In addition to the motion bots, a dedicated Arduino scale was designed that utilized a load cell and an HX711 amplifier and reads serial data from the Arduino based digital scale and sends the current mass as a serial message to `/weight`. This node parses the weight readings from `/dev/ttyACM0` and provides real-time feedback to the `SimplePourBot`, which uses the stream to stop the slow tilting motion once the 100 g threshold is reached.

## 2.6 End-Effector Cup Design for Stable Grasping

To ensure a robust grasp during fast motions and circular pouring trajectories, we designed a custom cup tailored to the Kinova Gen3 Lite gripper. We first measured the inner width and finger travel of the gripper and matched the cup diameter so that a fully closed grasp neither crushes the cup nor leaves excessive slack.

Using Autodesk Fusion, we created three surface variants: a smooth-walled cup, a cup with vertical ribs, and a cup with a diamond lattice texture. Each design was 3D printed and tested under real pouring motions to evaluate slip, vibration, and the risk of losing the grasp.

Experiments showed that the smooth cup was prone to slipping, and the vertically ribbed cup reduced slip but still allowed small shifts during dynamic motions. The diamond-patterned cup provided the most stable grasp, maximizing contact area and friction between the gripper pads and the cup surface. As a result, we selected the diamond-patterned design for all final gravimetric and barista pouring trials. Figure 4 shows the CAD models and the corresponding 3D printed prototypes.

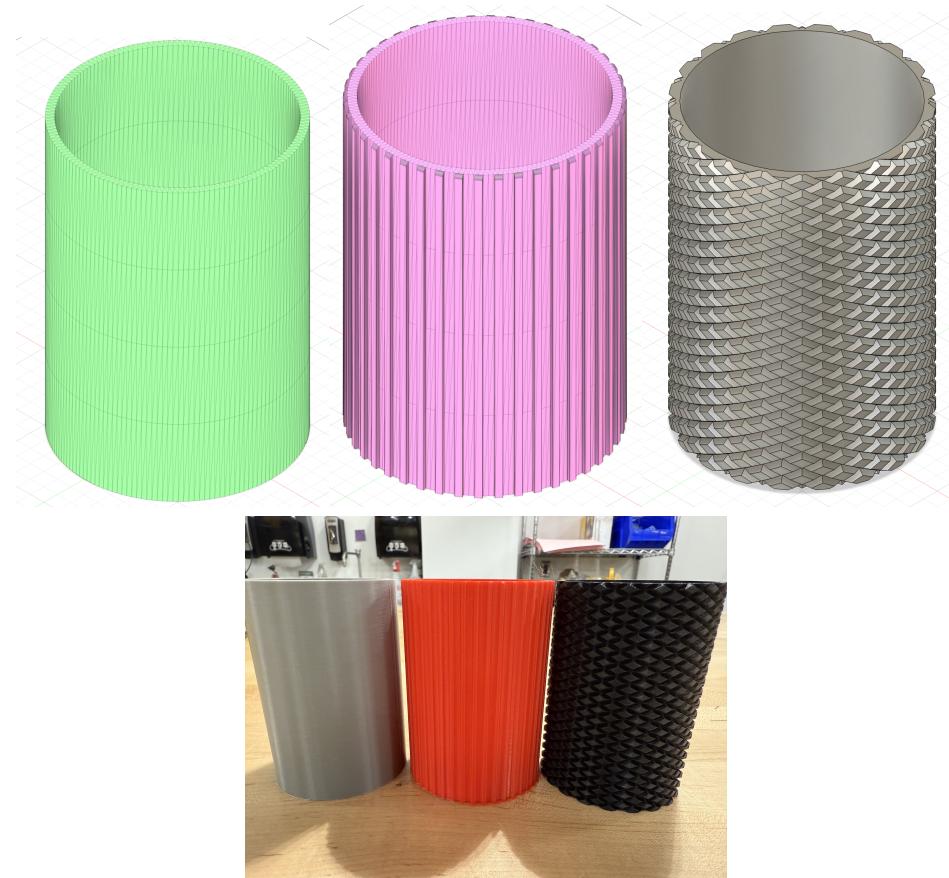


Figure 4: Custom cup designs and prototypes.

Top: CAD models of the smooth, ribbed, and diamond patterned cups.

Bottom: 3D printed prototypes used to evaluate grasp stability. The diamond patterned cup provided the most robust grasp and was selected for all final trials.

## 3 Human Integration

The system is designed for high-risk or sterile environments where physical contact is undesirable. The human operator interacts via a **non-contact visual interface**.

As illustrated in Figure 5, the interaction follows a safety protocol:

1. **Command Initiation:** Operator displays a hand gesture.
2. **Verification:** The node filters noise by requiring the signal to persist for 1 second.
3. **Execution & Lockout:** Once triggered, the system enters a **BUSY** state, ignoring all inputs until the task is complete.

## 4 Challenges

One significant challenge we faced was that YOLO could not run inside the default Docker environment provided for the course. Many required dependencies for ultralytics and OpenCV were missing, and direct installation attempts resulted in version conflicts. To resolve this, we created an extended Dockerfile based on the course image and manually added the necessary system libraries and Python packages. After rebuilding the container, we were finally able to run YOLO detection both offline and through the camera.

Another practical challenge involved the behavior-triggering mechanism for gesture control. Since the gesture topic continuously publishes the same detection result frame-by-frame, the robot would repeatedly execute the same action unless we introduced a trigger condition. To address this, we implemented a state-change mechanism: the robot only executes a motion when the detected gesture changes relative to the previous frame. This ensured stable and predictable control when integrating the vision system with the Kinova Gen3 arm.

A third challenge arose in the pouring dynamics and accuracy. In our initial implementation, the arm moved directly from the upright pose to a 90° tilt in a single motion. This caused a sudden surge of rice, often spilling outside the bowl and overshooting the 100 g target by a large margin, with final readings frequently in the 110–120 g range. We addressed this by redesigning the motion as a multi-step slow tilt. The cup orientation is interpolated between the start and end quaternions over 20 small steps while monitoring the weight. This modification made the motion visibly smoother and reduced the typical overshoot, so that the final mass after triggering the 100 g threshold stop was around 103 g.

## 5 Future Work

If provided more time, we would extend the gravimetric closed loop control that we implemented for the simple pour to the barista style circular pour. In the simple mode, a load cell, amplifier, and Arduino based scale already stream weight to ROS, and the robot stops pouring and returns the cup once the measured mass crosses a 100 g threshold while executing a 20 step, slowly tilting motion. However, in the barista pour mode the system still runs open loop. The arm tilts the cup once to a fixed angle and follows a circular trajectory without using the scale feedback, which can lead to over or under pouring. As future work, we would integrate the same weight feedback into the barista pour, modulate the circular path and tilt angle based on the real time mass reading, and redesign the motion as a multi step trajectory so that the circular pour remains smooth while stopping close to the desired target mass.

## 6 Takeaways

- **Concurrency in Robotics:** We learned that for any real-time interactive robot, single-threaded execution is insufficient. Mastering ROS 2 `Executors` and `CallbackGroups` is essential for creating responsive systems.
- **Parametric Trajectories:** We learned that simple point-to-point motion (PTP) is often unnatural for manipulation tasks. Generating trajectories using mathematical functions

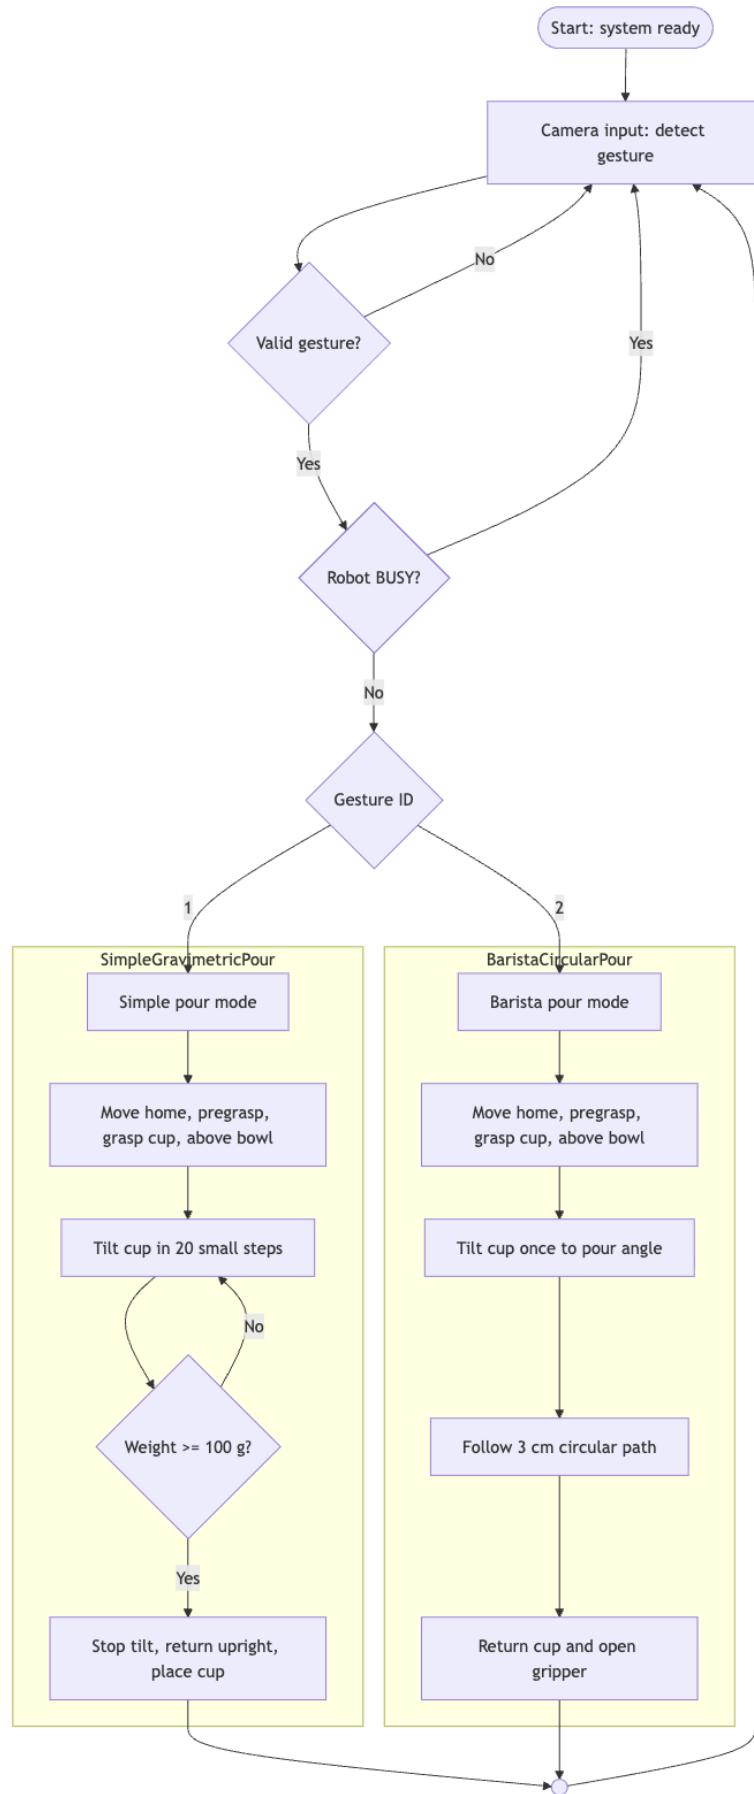


Figure 5: Flowchart of the Human-Robot Interaction Logic.

(like our circular pour) yields much smoother and more effective physical results than hard-coded waypoints.

- **Feedback and End-Effector Design:** We learned that achieving precise pouring requires both closed loop feedback and careful end effector design. The gravimetric pour showed how real time weight measurements reduce overshoot, and the custom diamond patterned cup significantly improved grasp stability during dynamic motions.

## A Code Appendix

### A.1 Gesture Detection Node

This node serves as the perception front-end. It performs real-time gesture recognition and publishes the resulting gesture label as a string message on the `/gesture` topic

Listing 1: Detect and Publish

```

1  #!/usr/bin/env python3
2  # predict_pub.py - debounced publisher that only republishes when label
3  # changes
4  import os
5  import sys
6  import time
7  import cv2
8  from ultralytics import YOLO
9
10 # ROS2 imports
11 import rclpy
12 from rclpy.node import Node
13 from std_msgs.msg import String
14
15 # -----
16 # Config (can override via env)
17 #
18 DEFAULT_MODEL = r"weights/best.pt"
19 MODEL_PATH = os.environ.get("YOLO_MODEL_PATH", DEFAULT_MODEL)
20 YOLO_VIDEO = os.environ.get("YOLO_TEST_VIDEO", "") # if set, use video
21     file
22 CAMERA_INDEX = int(os.environ.get("YOLO_VIDEO_INDEX", "0"))
23 CONF_THRESHOLD = float(os.environ.get("YOLO_CONF", "0.5"))
24 ROI_WIDTH = int(os.environ.get("YOLO_ROI_W", "200"))
25 ROI_HEIGHT = int(os.environ.get("YOLO_ROI_H", "200"))
26 DEBOUNCE_FRAMES = int(os.environ.get("YOLO_DEBOUNCE", "3"))
27 CLEAR_FRAMES = int(os.environ.get("YOLO_CLEAR_FRAMES", "8")) # consecutive "no-detection" frames to clear last_published_label
28 # Optional pause (not required with this method)
29 PAUSE_SEC = float(os.environ.get("YOLO_PAUSE_SEC", "0.0"))
30 #
31
32 def load_model(path):
33     try:
34         print(f"[INFO] Loading YOLO model from: {path}")
35         model = YOLO(path)
36         return model
37     except Exception as e:
38         print(f"[ERROR] Could not load model: {e}", file=sys.stderr)
39         raise
40
41 class YoloPublisher(Node):
42     def __init__(self, model, video_source, camera_index):
43         super().__init__('yolo_publisher')
44         self.pub = self.create_publisher(String, 'gesture', 10)
45         self.model = model
46         self.cap = None
47         if video_source:

```

```

47         self.get_logger().info(f"Opening video file: {video_source}")
48         self.cap = cv2.VideoCapture(video_source)
49     else:
50         self.get_logger().info(f"Opening camera index: {camera_index}")
51         self.cap = cv2.VideoCapture(camera_index)
52
53     if not self.cap or not self.cap.isOpened():
54         self.get_logger().error("Unable to open video/camera source")
55         raise RuntimeError("Unable to open video/camera source.")
56
57     # state for debounce and change-only logic
58     self.last_seen_label = None           # last label seen (for
59     debounce)
60     self.debounce_counter = 0
61
62     self.last_published_label = None      # last label we actually
63     published
64     self.clear_none_counter = 0          # counts consecutive frames
65     with no detection (to allow re-publish same label)
66
67     def busy_check_and_publish(self, label, conf):
68         # -----
69         # Called when we have a debounced label candidate.
70         # Publish only if label != last_published_label.
71         # After publishing, reset clear counter.
72         # -----
73         if label is None:
74             return False
75
76         # If label equals last published, ignore (require clear or
77         # different label)
78         if self.last_published_label == label:
79             self.get_logger().info(f"Label '{label}' equals last published; ignoring until change/clear.")
80             return False
81
82         # Publish now
83         msg = String()
84         msg.data = label
85         self.pub.publish(msg)
86         self.get_logger().info(f"Published gesture: {label} (conf={conf:.2f})")
87
88         # update state
89         self.last_published_label = label
90         self.clear_none_counter = 0
91         return True
92
93     def process_frame_and_publish(self):
94         ret, frame = self.cap.read()
95         if not ret:
96             # if video ended, rewind (useful for testing with short
97             # videos)
98             try:
99                 self.cap.set(cv2.CAP_PROP_POS_FRAMES, 0)
100             except Exception:

```

```

95         pass
96     return
97
98     h_frame, w_frame = frame.shape[:2]
99     center_x = int(w_frame * 0.75)
100    center_y = int(h_frame * 0.25)
101    roi_x1 = center_x - (ROI_WIDTH // 2)
102    roi_y1 = center_y - (ROI_HEIGHT // 2)
103    roi_x1 = max(0, min(roi_x1, w_frame - ROI_WIDTH))
104    roi_y1 = max(0, min(roi_y1, h_frame - ROI_HEIGHT))
105    roi_x2 = roi_x1 + ROI_WIDTH
106    roi_y2 = roi_y1 + ROI_HEIGHT
107
108    # draw guide box
109    cv2.rectangle(frame, (roi_x1, roi_y1), (roi_x2, roi_y2), (255,
110        0, 0), 2)
111    cv2.putText(frame, "Put Hand Here", (roi_x1, roi_y1 - 10),
112                cv2.FONT_HERSHEY_SIMPLEX, 0.6, (255, 0, 0), 2)
113
114    roi_frame = frame[roi_y1:roi_y2, roi_x1:roi_x2]
115
116    # Run inference (single ROI)
117    try:
118        results = self.model(roi_frame, stream=True, verbose=False)
119    except Exception as e:
120        self.get_logger().error(f"Inference error: {e}")
121        return
122
123    # Get best detection (label + confidence) or None
124    best_label = None
125    best_conf = 0.0
126    for result in results:
127        for box in result.bboxes:
128            conf = float(box.conf[0].item())
129            if conf < CONF_THRESHOLD:
130                continue
131            cls_idx = int(box.cls[0].item())
132            name = self.model.names.get(cls_idx, str(cls_idx)) if
133                hasattr(self.model, 'names') else str(cls_idx)
134            if conf > best_conf:
135                best_conf = conf
136                best_label = name
137
138    # --- Clear logic: if no detection in this frame, increment
139    # clear counter
140    if best_label is None:
141        self.clear_none_counter += 1
142    else:
143        self.clear_none_counter = 0
144
145    # --- Debounce logic: require DEBOUNCE_FRAMES consecutive same
146    # label to consider it real
147    if best_label is None:
148        # reset last seen label and debounce counter
        self.last_seen_label = None
        self.debounce_counter = 0
    else:
        if best_label == self.last_seen_label:

```

```

149         self.debounce_counter += 1
150     else:
151         self.last_seen_label = best_label
152         self.debounce_counter = 1
153
154     # When debounced candidate ready:
155     if self.debounce_counter >= DEBOUNCE_FRAMES:
156         # Only publish if label differs from last published OR if
157         # last_published_label is None
158         published = self.busy_check_and_publish(self.
159             last_seen_label, best_conf)
160         # optional small hard pause after publish
161         if published and PAUSE_SEC > 0:
162             self.get_logger().info(f"Hard\u00a0pause\u00a0for\u00a0{PAUSE_SEC}\u00a0
163             seconds\u00a0after\u00a0publish.")
164             time.sleep(PAUSE_SEC)
165         # reset debounce so we don't immediately re-evaluate same
166         # frames
167         self.debounce_counter = 0
168         self.last_seen_label = None
169
170     # If last published label exists and we've seen enough
171     # consecutive empty frames, clear it so same label can re-
172     # trigger
173     if self.last_published_label is not None and self.
174         clear_none_counter >= CLEAR_FRAMES:
175         self.get_logger().info("Clear\u00a0threshold\u00a0reached:\u00a0allowing\u00a0
176         re-trigger\u00a0of\u00a0last\u00a0published\u00a0label.")
177         self.last_published_label = None
178         self.clear_none_counter = 0
179
180     # show frame (optional)
181     cv2.imshow('ROI\u00a0Hand\u00a0Gesture\u00a0Detection', frame)
182     if cv2.waitKey(1) & 0xFF == ord('q'):
183         raise KeyboardInterrupt()
184
185     def shutdown(self):
186         try:
187             self.cap.release()
188         except Exception:
189             pass
190         cv2.destroyAllWindows()
191
192     def main(args=None):
193         # load model
194         model = load_model(MODEL_PATH)
195
196         # choose video or camera
197         video_src = YOLO_VIDEO if YOLO_VIDEO else None
198
199         # init ROS
200         rclpy.init(args=args)
201         node = YoloPublisher(model=model, video_source=video_src,
202             camera_index=CAMERA_INDEX)
203
204         try:
205             node.get_logger().info("Starting\u00a0detection\u00a0loop.\u00a0Press\u00a0'q'\u00a0in\u00a0
206             window\u00a0to\u00a0quit.")

```

```
197     while rclpy.ok():
198         node.process_frame_and_publish()
199         rclpy.spin_once(node, timeout_sec=0)
200     except KeyboardInterrupt:
201         node.get_logger().info("Keyboard interrupt, shutting down.")
202     finally:
203         node.shutdown()
204         node.destroy_node()
205         rclpy.shutdown()
206
207 if __name__ == "__main__":
208     main()
```

## A.2 Gesture Control (Supervisor Node)

This node handles the state machine and multi-threading logic.

Listing 2: Gesture Control Logic

```
1  #!/usr/bin/env python3
2
3
4  import rclpy
5  import threading
6  from rclpy.node import Node
7  from rclpy.executors import MultiThreadedExecutor
8  from std_msgs.msg import String
9  import time
10
11 # CRITICAL: These files must NOT have their own main() and rclpy.init()
12 # blocks.
13 # They are imported here as service classes (Nodes).
14 from .coffee_pourer import SimpleCoffeeBot as SimplePourBot
15 from .coffee_pourer_v2 import SimpleCoffeeBot as BaristaPourBot
16
17 class GestureControl(Node):
18     # -----
19     # Central node that manages robot action nodes and listens for
20     # commands
21     # From the YOLO vision publisher. Uses MultiThreadedExecutor for
22     # stability.
23     # -----
24
25     def __init__(self, action_node_executor):
26         super().__init__("gesture_control_node")
27
28         # --- State Management ---
29         self.busy = False
30         self.executor = action_node_executor
31
32         # --- Action Node Instantiation (Created only ONCE) ---
33         # These objects are the actual robot control nodes that talk to
34         # MoveIt2.
35         self.simple_bot = SimplePourBot()
36         self.barista_bot = BaristaPourBot()
37
38         # Add the action nodes to the shared Executor so they can
39         # process ROS topics
40         self.executor.add_node(self.simple_bot)
41         self.executor.add_node(self.barista_bot)
42
43         # --- ROS Subscription ---
44         # Ensure this topic name matches the one published by your YOLO
45         # node
46         self.sub = self.create_subscription(String, "gesture", self.
47             gesture_callback, 10)
48
49         self.get_logger().info("GestureControl and Robot Managers are
50             ready.")
51         self.get_logger().info("Listening to /gesture topic for
52             commands (1 or 2)...")
```

```

45
46     def gesture_callback(self, msg):
47         """Processes incoming gesture commands."""
48         label = msg.data.strip()
49         self.get_logger().info(f"Received gesture: {label}")
50
51         if self.busy:
52             self.get_logger().warn("Robot busy - ignoring gesture.")
53             return
54
55         # --- Modified Logic: Check for '1' or '2' ---
56         if label == "1":
57             # Pass the run_my_code method of the pre-initialized bot
58             # instance
59             self.start_thread(self.simple_bot.run_my_code)
60             self.get_logger().info("Triggering: SimplePour(V1)")
61
62         elif label == "2":
63             self.start_thread(self.barista_bot.run_my_code)
64             self.get_logger().info("Triggering: BaristaPour(V2)")
65
66         else:
67             self.get_logger().warn(f"Unknown gesture: {label}")
68
69         # =====
70         # Thread management
71         # =====
72     def start_thread(self, target_fn):
73         """Starts the robot action in a separate thread to prevent ROS
74             callbacks from blocking."""
75         thread = threading.Thread(target=self.safe_run, args=(target_fn,
76             ), daemon=True)
77         thread.start()
78
79     def safe_run(self, target_fn):
80         """Wrapper function to manage the robot's busy state."""
81         self.busy = True
82         self.get_logger().info(f"Starting action: {target_fn.__name__}")
83
84         try:
85             # Execute the actual motion function
86             target_fn()
87         except Exception as e:
88             self.get_logger().error(f"Error executing motion: {e}")
89         finally:
90             self.busy = False
91             self.get_logger().info("Action finished. Ready for next
92             command.")
93
94     def main(args=None):
95         rclpy.init(args=args)
96
97         # 1. Create MultiThreaded Executor to handle callbacks and MoveIt2
98             # actions concurrently
99         executor = MultiThreadedExecutor(num_threads=4)
100
101         # 2. Create the main control node (passing the executor)

```

```
97     control_node = GestureControl(executor)
98
99     # 3. Add the control node to the executor
100    executor.add_node(control_node)
101
102    try:
103        # 4. Start the ROS event loop
104        executor.spin()
105    except KeyboardInterrupt:
106        pass
107    except Exception as e:
108        control_node.get_logger().error(f"Executor error: {e}")
109    finally:
110        rclpy.shutdown()
111
112
113 if __name__ == "__main__":
114     main()
```

### A.3 Simple Pour Node (Static)

This node handles the basic Pick-and-Pour logic without trajectory modification.

Listing 3: Simple Coffee Pouring

```
1 #!/usr/bin/env python3
2
3
4 import time
5 import threading
6 import math
7 import serial
8 import rclpy
9 from rclpy.node import Node
10 from rclpy.executors import MultiThreadedExecutor
11 from geometry_msgs.msg import Pose
12 from pymoveit2 import MoveIt2
13 from pymoveit2.gripper_interface import GripperInterface
14
15
16 class SimpleCoffeeBot(Node):
17     def __init__(self):
18         super().__init__("simple_coffee_node")
19
20         # --- SERIAL SETUP ---
21         # Initialize the serial connection when the node starts
22         self.serial_port_name = "/dev/ttyACM0" # Check this device
23             name
24         self.baud_rate = 9600
25         self.ser = None
26
27         try:
28             self.ser = serial.Serial(self.serial_port_name, self.
29                 baud_rate, timeout=1)
30             self.get_logger().info(f"Connected to serial device on {self.
31                 serial_port_name}")
32         except serial.SerialException as e:
33             self.get_logger().error(f"Failed to connect to serial: {e}")
34
35         # --- ROBOT SETUP ---
36         self.arm = MoveIt2(
37             node=self,
38             joint_names=["joint_1", "joint_2", "joint_3", "joint_4", "joint_5", "joint_6"],
39             base_link_name="base_link",
40             end_effector_name="end_effector_link",
41             group_name="arm",
42         )
43
44         # Gripper setup
45         self.gripper = GripperInterface(
46             node=self,
47             gripper_joint_names=["right_finger_bottom_joint"],
48             open_gripper_joint_positions=[0.01],
49             closed_gripper_joint_positions=[0.8],
50             gripper_group_name="gripper",
```

```

48         gripper_command_action_name="/
49             gen3_lite_2f_gripper_controller/gripper_cmd",
50             ignore_new_calls_while_executing=False,
51     )
52
53     self.get_logger().info("Robot is ready!")
54
55 # -----
56 # Quaternion helpers
57 # -----
58 def lerp_quat(self, q0, q1, t: float):
59     """Linear interpolation between two quaternions followed by normalization."""
60     q = [q0[i] + t * (q1[i] - q0[i]) for i in range(4)]
61     norm = math.sqrt(sum(v * v for v in q))
62     if norm == 0.0:
63         return q0[:] # Fallback to start quaternion if something goes wrong
64     return [v / norm for v in q]
65
66 def tilt_slowly(self, x, y, z, q_start, q_end, steps: int = 10):
67     """Gradually tilt the cup from q_start to q_end in multiple small steps.
68     This makes the pouring motion look slow and smooth instead of sudden.
69     """
70     for i in range(1, steps + 1):
71         t = i / steps # interpolation factor from 0 to 1
72         q = self.lerp_quat(q_start, q_end, t)
73         self.move_to(x, y, z, q)
74
75 # -----
76 # Main sequence
77 # -----
78 def run_my_code(self):
79     # --- COORDINATES ---
80     cup_x = 0.40
81     cup_y = -0.20
82     cup_z = 0.0001
83
84     pour_x = 0.40
85     pour_y = 0.20
86     pour_z = 0.22
87
88     # Orientation: cup held straight
89     orient_straight = [0.0, 0.7071, 0.0, 0.7071]
90
91     # Orientation: tilted for pouring (from RViz calibration)
92     orient_tilted = [0.5373, 0.4869, 0.4869, 0.4869] # ~115 degrees
93
94     # -----
95     # STEP 1: Go to a safe home pose and open gripper
96     # -----
97     self.get_logger().info("Step 1: Go to home and open gripper")
98     self.arm.move_to_configuration(
99         joint_positions=[0.0, -0.28, 1.3, 1.0, -1.0, 0.0]

```

```
100
101     )
102     self.arm.wait_until_executed()
103
104     self.gripper.open()
105     # If you want to be safer, you can also call self.gripper.
106     # wait_until_executed()
107     time.sleep(1.0)
108
109     # -----
110     # STEP 2: Pick up the cup
111     # -----
112     self.get_logger().info("Step 2: Picking up the cup")
113
114     # Move above the cup
115     self.move_to(cup_x, cup_y, cup_z + 0.15, orient_straight)
116     # Move down to the cup
117     self.move_to(cup_x, cup_y, cup_z, orient_straight)
118
119     # Close gripper to grab the cup
120     self.gripper.close()
121     # Optionally: self.gripper.wait_until_executed()
122     time.sleep(1.0)
123
124     # Lift the cup up
125     self.move_to(cup_x, cup_y, cup_z + 0.15, orient_straight)
126
127     # -----
128     # STEP 3: Move to the pouring position
129     # -----
130     self.get_logger().info("Step 3: Moving to the pouring position")
131     self.move_to(pour_x, pour_y, pour_z, orient_straight)
132
133     # -----
134     # STEP 4: Pour (tilt slowly and then wait for serial)
135     # -----
136     self.get_logger().info("Step 4: Pouring slowly...")
137     # Slowly tilt from straight orientation to tilted orientation
138     self.tilt_slowly(
139         pour_x,
140         pour_y,
141         pour_z,
142         orient_straight,
143         orient_tilted,
144         steps=20,
145     )
146
147     # Wait until the sensor value exceeds a threshold (e.g., weight
148     # > 100)
149     if self.ser and self.ser.is_open:
150         self.get_logger().info("Reading serial data... waiting for
151         value > 100")
152         # Clear any old data from the buffer so we don't read stale
153         # values
154         self.ser.reset_input_buffer()
155
156         while True:
157             try:
```

```

153         # Read a line, decode from bytes to string, strip
154         # whitespace
155         line = self.ser.readline().decode("utf-8", errors="ignore").strip()
156
157     if line:
158         try:
159             value = float(line)
160             self.get_logger().info(f"Sensor Value: {value}")
161             if value > 100:
162                 self.get_logger().info("Target reached!
163                             Stopping pour.")
164                 break
165         except ValueError:
166             # Ignore non-numeric messages, but log them
167             # for debugging
168             self.get_logger().warn(f"Received non-
169             numeric data: {line}")
170
171     except Exception as e:
172         self.get_logger().error(f"Serial read error: {e}")
173         break
174
175 else:
176     self.get_logger().warn("Serial not connected! Defaulting to
177                             2 second wait.")
178     time.sleep(2.0)
179
180 # -----
181 # STEP 5: Straighten the cup again
182 # -----
183 self.get_logger().info("Step 5: Straightening the cup")
184 self.move_to(pour_x, pour_y, pour_z, orient_straight)
185
186 # -----
187 # STEP 6: Return the cup to its original position
188 # -----
189 self.get_logger().info("Step 6: Returning the cup")
190 # Move above the original cup position
191 self.move_to(cup_x, cup_y, cup_z + 0.15, orient_straight)
192 # Move down to place the cup
193 self.move_to(cup_x, cup_y, cup_z, orient_straight)
194
195 # Open gripper to release the cup
196 self.gripper.open()
197 # Optionally: self.gripper.wait_until_executed()
198 time.sleep(1.0)
199
200 # Move arm up again
201 self.move_to(cup_x, cup_y, cup_z + 0.15, orient_straight)
202
203 # Go back to a neutral joint configuration
204 self.get_logger().info("Done! Moving back to neutral
205                         configuration.")
206 self.arm.move_to_configuration(
207     joint_positions=[0.0, 0.0, 0.0, 0.0, 0.0, 0.0]
208 )
209 self.arm.wait_until_executed()

```

```
203
204     # -----
205     # Helper: move to a Cartesian pose
206     #
207     def move_to(self, x, y, z, orientation):
208         """Plan and execute a Cartesian move to the given pose."""
209         pose = Pose()
210         pose.position.x = x
211         pose.position.y = y
212         pose.position.z = z
213         pose.orientation.x = orientation[0]
214         pose.orientation.y = orientation[1]
215         pose.orientation.z = orientation[2]
216         pose.orientation.w = orientation[3]
217
218         plan = self.arm.plan(pose=pose, cartesian=True)
219         if plan:
220             self.arm.execute(plan)
221             self.arm.wait_until_executed()
222         else:
223             self.get_logger().error("Could not find a path!")
224
225
226     # -----
227     # Standard ROS2 boilerplate
228     #
229     def main(args=None):
230         rclpy.init(args=args)
231         node = SimpleCoffeeBot()
232
233         executor = MultiThreadedExecutor(num_threads=2)
234         executor.add_node(node)
235         thread = threading.Thread(target=executor.spin, daemon=True)
236         thread.start()
237
238         try:
239             # Small delay to let everything start up
240             time.sleep(2.0)
241             node.run_my_code()
242         except KeyboardInterrupt:
243             pass
244         finally:
245             # Close serial port cleanly on exit
246             if node.ser and node.ser.is_open:
247                 node.ser.close()
248             rclpy.shutdown()
249             thread.join()
```

## A.4 Barista Pour Node (Parametric)

This node implements the circular path planning algorithm.

Listing 4: Barista Circular Path

```

1  #!/usr/bin/env python3
2
3
4  import time
5  import threading
6  import rclpy
7  import math
8  from rclpy.node import Node
9  from rclpy.executors import MultiThreadedExecutor
10 from geometry_msgs.msg import Pose
11 from pymoveit2 import MoveIt2
12 from pymoveit2.gripper_interface import GripperInterface
13
14
15 class SimpleCoffeeBot(Node):
16     def __init__(self):
17         super().__init__("simple_coffee_node")
18
19         # --- ROBOT SETUP (DON'T TOUCH THIS) ---
20         self.arm = MoveIt2(
21             node=self,
22             joint_names=["joint_1", "joint_2", "joint_3", "joint_4",
23                         "joint_5", "joint_6"],
24             base_link_name="base_link",
25             end_effector_name="end_effector_link",
26             group_name="arm",
27         )
28
29         # Setup the Gripper (The fingers)
30         self.gripper = GripperInterface(
31             node=self,
32             gripper_joint_names=["right_finger_bottom_joint"],
33             open_gripper_joint_positions=[0.01],      # Open
34             closed_gripper_joint_positions=[0.8],    # Closed
35             gripper_group_name="gripper",
36             gripper_command_action_name="/
37                 gen3_lite_2f_gripper_controller/gripper_cmd",
38             ignore_new_calls_while_executing=False,
39         )
40         self.get_logger().info("Robot is ready!")
41
42         # =====
43         # YOUR CODE GOES HERE (READ THIS PART)
44         # =====
45         def run_my_code(self):
46
47             # --- COORDINATES (Change these numbers!) ---
48             cup_x = 0.40
49             cup_y = -0.20
50             cup_z = 0.01
51
52             pour_x = 0.40
53             pour_y = 0.20

```

```

52     pour_z = 0.22
53
54     # Orientation: Normal (Holding cup straight)
55     orient_straight = [0.0, 0.7071, 0.0, 0.7071]
56
57     # Orientation: Tilted (Pouring)
58     orient_tilted = [0.5, 0.5, 0.5, 0.9]
59
60     # -----
61     # STEP 1: Go Home and Open Hand
62     # -----
63     self.get_logger().info("Step 1: Start")
64     # Move joints to a safe starting position
65     self.arm.move_to_configuration(joint_positions=[0.0, -0.28,
66                                         1.3, 1.0, -1.0, 0.0])
67     self.arm.wait_until_executed()
68
68     self.gripper.open()
69     time.sleep(1.0)
70
71     # -----
72     # STEP 2: Pick up the Cup
73     # -----
74     self.get_logger().info("Step 2: Picking up cup")
75
76     # A. Go above the cup (Hover)
77     self.move_to(cup_x, cup_y, cup_z + 0.15, orient_straight)
78
79     # B. Go down to the cup
80     self.move_to(cup_x, cup_y, cup_z, orient_straight)
81
82     # C. Grab it
83     self.gripper.close()
84     time.sleep(1.0) # Wait for grip
85
86     # D. Lift it up
87     self.move_to(cup_x, cup_y, cup_z + 0.15, orient_straight)
88
89     # -----
90     # STEP 3: Go to Pouring Spot
91     # -----
92     self.get_logger().info("Step 3: Moving to pour spot")
93     # Move to the pour location, still holding it straight
94     self.move_to(pour_x, pour_y, pour_z, orient_straight)
95
96     # -----
97     # STEP 4: POUR! (Tilt the robot) old version
98     # -----
99     # self.get_logger().info("Step 4: Pouring...")
100    # # Stay in same spot, change orientation to TILTED
101    # self.move_to(pour_x, pour_y, pour_z, orient_tilted)
102
103    # time.sleep(2.0) # Wait for coffee to pour out...
104
105    # -----
106    # STEP 4: The "Barista" Circular Pour
107    # -----
108    self.get_logger().info("Step 4: Barista Mode Initiated")

```

```

109
110     # 1. Move to the START of the circle (Angle 0)
111     # We start at x + radius, so we don't crash into the side
112     # immediately
113     radius = 0.03 # 3cm radius circle
114     start_x = pour_x + radius
115     start_y = pour_y
116
117     # Tilt the cup first at the starting point
118     self.move_to(start_x, start_y, pour_z, orient_tilted)
119
120     # 2. Perform the Circle
121     # We divide the circle into 20 steps
122     steps = 20
123
124     for i in range(steps):
125         # Calculate the angle for this step (0 to 360 degrees)
126         angle = (i / steps) * 2.0 * math.pi
127
128         # Calculate the new X and Y based on the circle formula
129         # Note: We keep Z and Orientation CONSTANT
130         next_x = pour_x + (radius * math.cos(angle))
131         next_y = pour_y + (radius * math.sin(angle))
132
133         self.get_logger().info(f"Pouring\u00b7at\u00b7angle:\u00b7{int(math.
134             degrees(angle))}")
135
136         # Move to the next point on the circle
137         self.move_to(next_x, next_y, pour_z, orient_tilted)
138
139         # -----
140         # STEP 5: Stop Pouring
141         # -----
142         self.get_logger().info("Step\u20225:\u2022Straightening\u00b7up")
143         # Change orientation back to STRAIGHT
144         self.move_to(pour_x, pour_y, pour_z, orient_straight)
145
146         # -----
147         # STEP 6: Put Cup Back
148         # -----
149         self.get_logger().info("Step\u20226:\u2022Returning\u00b7cup")
150
151         # A. Go above the original spot
152         self.move_to(cup_x, cup_y, cup_z + 0.15, orient_straight)
153
154         # B. Go down
155         self.move_to(cup_x, cup_y, cup_z, orient_straight)
156
157         # C. Let go
158         self.gripper.open()
159         time.sleep(1.0)
160
161         # D. Move away (Up)
162         self.move_to(cup_x, cup_y, cup_z + 0.15, orient_straight)
163
164         self.get_logger().info("Done!")

```

```
165     self.arm.move_to_configuration(joint_positions=[0.0, 0.0, 0.0,
166                                         0.0, 0.0, 0.0])
167     self.arm.wait_until_executed()
168
169     self.get_logger().info("Done!")
170
171 # =====
172 # HELPER FUNCTION (Makes moving easier)
173 # =====
174 def move_to(self, x, y, z, orientation):
175     # This function handles all the boring math to create a "Pose"
176     pose = Pose()
177     pose.position.x = x
178     pose.position.y = y
179     pose.position.z = z
180     pose.orientation.x = orientation[0]
181     pose.orientation.y = orientation[1]
182     pose.orientation.z = orientation[2]
183     pose.orientation.w = orientation[3]
184
185     # Tell the robot to plan and execute the path
186     plan = self.arm.plan(pose=pose, cartesian=True)
187     if plan:
188         self.arm.execute(plan)
189         self.arm.wait_until_executed()
190     else:
191         self.get_logger().error("Could not find a path!")
192
193 # -----
194 # Standard Boilerplate (Just copy this)
195 # -----
196 def main(args=None):
197     rclpy.init(args=args)
198     node = SimpleCoffeeBot()
199
200     # Run the node in a background thread
201     executor = MultiThreadedExecutor(num_threads=2)
202     executor.add_node(node)
203     thread = threading.Thread(target=executor.spin, daemon=True)
204     thread.start()
205
206     try:
207         time.sleep(2.0) # Wait for startup
208         node.run_my_code() # <-- This runs your code above
209     except KeyboardInterrupt:
210         pass
211     finally:
212         rclpy.shutdown()
213         thread.join()
```

## A.5 Arduino Code for Scale

This standalone script runs on Arduino, using a load cell and an HX711 amplifier to output the mass to the serial monitor, which can then be read by the ROS 2 script.

Listing 5: Arduino Code for Scale

```

1 #include <Arduino.h>
2 #include "HX711.h"
3
4 // circuit wiring
5 const int LOADCELL_DOUT_PIN = 2;
6 const int LOADCELL_SCK_PIN = 3;
7 const int BUTTON_PIN = 7;
8
9
10 long offset = 566200;
11 const float scale_factor = 420.58; // haha get it
12
13 HX711 scale;
14
15 /**
16 * @brief Performs a button-triggered tare operation.
17 * It reads the HX711 multiple times over a 1-second period
18 * and calculates the average raw reading to set the new offset.
19 * @return The new calculated average offset.
20 */
21 long button_tare() {
22     const int TARE_DURATION_MS = 1000; // Length of taring sequence (ms)
23     const int SAMPLE_DELAY_MS = 50; // Delay between samples
24     long sumOfReadings = 0; // Initialize running total
25     int count = 0;
26     unsigned long startTime = millis();
27
28     Serial.println("\n---\u2022Taring\u2022in\u2022progress\u2022(1\u2022second)\u2022---");
29
30     // Read the scale for the specified duration
31     while (millis() - startTime < TARE_DURATION_MS) {
32         if (scale.is_ready()) {
33             sumOfReadings += scale.read();
34             count++;
35         }
36         delay(SAMPLE_DELAY_MS);
37     }
38
39     // Calculate the average offset
40     if (count > 0) {
41         long newOffset = sumOfReadings / count;
42         Serial.print("Tare\u2022complete.\u2022Averaged\u2022");
43         Serial.print(count);
44         Serial.print("\u2022readings.\u2022New\u2022Offset:\u2022");
45         Serial.println(newOffset);
46         return newOffset;
47     } else {
48         Serial.println("Error:\u2022Scale\u2022not\u2022ready\u2022during\u2022tare,\u2022keeping\u2022old\u2022offset.\u2022");
49     }
50     return offset; // Return the current offset if no readings were
51     taken
52 }
```

```
51 }
52 }
53
54 void setup() {
55   Serial.begin(9600);
56   Serial.println("HX711 Raw Reading Demo");
57
58   // Initialize the scale with the DOUT and SCK pins
59   scale.begin(LOADCELL_DOUT_PIN, LOADCELL_SCK_PIN);
60   scale.tare();
61
62   // Pin Setup
63   pinMode(BUTTON_PIN, INPUT);
64 }
65
66 void loop() {
67   if (digitalRead(BUTTON_PIN) == HIGH) {
68     // Button is pressed, enter the taring state
69     offset = button_tare();
70   }
71
72   // Check if the scale is ready to read
73
74   if (scale.is_ready()) {
75     // scale.read() returns the raw 24-bit reading from the ADC
76     long reading = (scale.read() - offset) / scale_factor;
77     // Serial.print("Raw Reading: ");
78     Serial.println(reading);
79   } else {
80     Serial.println("HX711 not found.");
81   }
82
83   // A small delay to keep the serial monitor readable
84   delay(500);
85 }
```

## A.6 YOLO Dataset Configuration (data.yaml)

This configuration file serves as the bridge between the YOLO model and the training data, defining the file paths, number of classes (nc), and the human-readable class names (names).

Listing 6: YOLO Dataset Configuration (data.yaml)

```

1 # Dataset root directory (Absolute path recommended)
2 path: '/path/to/dataset/gesture'
3
4 # Train and Validation image directories (Relative to 'path')
5 train: images\train
6 val: images\val
7
8 # Number of classes
9 nc: 10
10
11 # Class names
12 names: ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']

```

## A.7 Computer Vision: Training Script

This script handles the transfer learning process, loading the pre-trained YOLOv8n weights and fine-tuning them on the hand gesture dataset.

Listing 7: YOLO Training Script

```

1 from ultralytics import YOLO
2 # Main entry point is required for Windows multiprocessing
3 if __name__ == '__main__':
4
5     # --- Configuration ---
6     # Path to your CORRECTED data.yaml file
7     data_yaml_path = r"data.yaml"
8
9     # Path to the pre-trained model (yolov8n.pt will be downloaded if
10       not found)
11     pre_model_path = "yolov8n.pt"
12
13     # Project name for saving results
14     project_name = "hand_gesture_train"
15
16     # Device setup: 0 for GPU, 'cpu' for CPU (Assuming CPU for
17       compatibility)
18     device_id = 'cpu'
19
20     # Training parameters
21     NUM_EPOCHS = 30
22     BATCH_SIZE = 16
23     NUM_WORKERS = 2
24     # -----
25
26     print("Loading model and transferring weights...")
27     # Initialize the model and load pre-trained weights
28     model = YOLO('yolov8n.yaml').load(pre_model_path)
29
30     print("Starting training...")
31     # Start training using the corrected configuration
32     results = model.train(

```

```
31     data=data_yaml_path ,
32     epochs=NUM_EPOCHS ,
33     imgsz=640 ,
34     batch=BATCH_SIZE ,
35     workers=NUM_WORKERS ,
36     name=project_name ,
37     device=device_id
38 )
39
40     print(f"Training completed. Best weights saved in runs/detect/{project_name}/weights/best.pt")
```